

# Class Challenge: Image Classification of COVID-19 X-rays

## Task 1 [Total points: 30]

### Setup

- This assignment involves the following packages: 'matplotlib', 'numpy', and 'sklearn'.
- If you are using conda, use the following commands to install the above packages:

```
conda install matplotlib
conda install numpy
conda install -c anaconda scikit-learn
```

- If you are using pip, use the following commands to install the above packages:

```
pip install matplotlib
pip install numpy
pip install sklearn
```

### Data

Please download the data using the following link: [COVID-19 \(https://drive.google.com/file/d/1Y88tgqpQ1Pjko\\_7rntcPowOJs\\_QNOrJ-/view\)](https://drive.google.com/file/d/1Y88tgqpQ1Pjko_7rntcPowOJs_QNOrJ-/view).

- After downloading 'Covid\_Data\_GradientCrescent.zip', unzip the file and you should see the following data structure:

```
--all
-----train
-----test
--two
-----train
-----test
```

- Put the 'all' folder, the 'two' folder and this python notebook in the **same directory** so that the following code can correctly locate the data.

## [20 points] Binary Classification: COVID-19 vs. Normal

In [1]:

```
import os

import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# os.environ['OMP_NUM_THREADS'] = '1'
# os.environ["CUDA_DEVICE_ORDER"]="PCI_BUS_ID"
# os.environ['CUDA_VISIBLE_DEVICES'] = '2'
tf.__version__
```

Out[1]:

'2.6.0'

In [2]:

```
from tensorflow.python.client import device_lib
print(device_lib.list_local_devices())
```

```
[name: "/device:CPU:0"
device_type: "CPU"
memory_limit: 268435456
locality {
}
incarnation: 7688805591285904748
, name: "/device:GPU:0"
device_type: "GPU"
memory_limit: 6252920832
locality {
  bus_id: 1
  links {
  }
}
incarnation: 1699347087621046735
physical_device_desc: "device: 0, name: NVIDIA GeForce RTX 2080 SUPE
R, pci bus id: 0000:01:00.0, compute capability: 7.5"
]
```

## Load Image Data

In [3]:

```
DATA_LIST = os.listdir('two/train')
DATASET_PATH = 'two/train'
TEST_DIR = 'two/test'
IMAGE_SIZE = (224, 224)
NUM_CLASSES = len(DATA_LIST)
BATCH_SIZE = 10 # try reducing batch size or freeze more layers if your GPU
runs out of memory
NUM_EPOCHS = 40
LEARNING_RATE = 0.0005 # start off with high rate first 0.001 and experiment wit
h reducing it gradually
```

## Generate Training and Validation Batches

In [4]:

```
train_datagen = ImageDataGenerator(rescale=1./255,rotation_range=50,featurewise_
center = True,
                                featurewise_std_normalization = True,width_sh
ift_range=0.2,
                                height_shift_range=0.2,shear_range=0.25,zoom_
range=0.1,
                                zca_whitening = True,channel_shift_range = 20
                                ,
                                horizontal_flip = True,vertical_flip = True,
                                validation_split = 0.2,fill_mode='constant')

train_batches = train_datagen.flow_from_directory(DATASET_PATH,target_size=IMAGE
_SIZE,
                                                shuffle=True,batch_size=BATCH_
SIZE,
                                                subset = "training",seed=42,
                                                class_mode="binary")

valid_batches = train_datagen.flow_from_directory(DATASET_PATH,target_size=IMAGE
_SIZE,
                                                shuffle=True,batch_size=BATCH_
SIZE,
                                                subset = "validation",seed=42,
                                                class_mode="binary")
```

Found 104 images belonging to 2 classes.

Found 26 images belonging to 2 classes.

C:\Users\Li\anaconda3\lib\site-packages\keras\_preprocessing\image\im  
age\_data\_generator.py:342: UserWarning: This ImageDataGenerator spec  
ifies `zca\_whitening` which overrides setting of `featurewise\_std\_nor  
malization`.

```
warnings.warn('This ImageDataGenerator specifies '
```

## [10 points] Build Model

Hint: Starting from a pre-trained model typically helps performance on a new task, e.g. starting with weights obtained by training on ImageNet.

In [5]:

```
from tensorflow.keras.applications import VGG16
from keras.layers.core import Flatten, Dense, Dropout, Lambda

vgg16 = VGG16(weights='imagenet', include_top=False, pooling = "None", classes = 2
, input_shape=(224, 224, 3))
vgg16.trainable = False

model = tf.keras.models.Sequential()
model.add(vgg16)
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(256, activation='relu'))
model.add(tf.keras.layers.Dense(1, activation='sigmoid'))
model.summary()

model.compile(loss='binary_crossentropy', optimizer=tf.keras.optimizers.Adam(learning_rate=LEARNING_RATE), metrics=['accuracy'])
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
vgg16 (Functional)	(None, 7, 7, 512)	14714688
-----		
flatten (Flatten)	(None, 25088)	0
-----		
dense (Dense)	(None, 256)	6422784
-----		
dense_1 (Dense)	(None, 1)	257
=====		
Total params: 21,137,729		
Trainable params: 6,423,041		
Non-trainable params: 14,714,688		

## [5 points] Train Model

In [6]:

```
#FIT MODEL
print(len(train_batches))
print(len(valid_batches))

STEP_SIZE_TRAIN=train_batches.n//train_batches.batch_size
STEP_SIZE_VALID=valid_batches.n//valid_batches.batch_size

import time
start = time.time()

history=model.fit_generator(train_batches, steps_per_epoch =STEP_SIZE_TRAIN, val
idation_data = valid_batches, validation_steps = STEP_SIZE_VALID, epochs= NUM_EP
OCHS)

end = time.time()
total_time = end - start
```

11

3

C:\Users\Li\anaconda3\lib\site-packages\keras\engine\training.py:197  
2: UserWarning: `Model.fit\_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

warnings.warn(`Model.fit\_generator` is deprecated and '  
C:\Users\Li\anaconda3\lib\site-packages\keras\_preprocessing\image\image\_data\_generator.py:720: UserWarning: This ImageDataGenerator specifies `featurewise\_center`, but it hasn't been fit on any training data. Fit it first by calling `.fit(numpy\_data)`.

warnings.warn('This ImageDataGenerator specifies '  
C:\Users\Li\anaconda3\lib\site-packages\keras\_preprocessing\image\image\_data\_generator.py:739: UserWarning: This ImageDataGenerator specifies `zca\_whitening`, but it hasn't been fit on any training data. Fit it first by calling `.fit(numpy\_data)`.

warnings.warn('This ImageDataGenerator specifies '

Epoch 1/40

10/10 [=====] - 6s 340ms/step - loss: 2.3690 - accuracy: 0.4149 - val\_loss: 0.7184 - val\_accuracy: 0.5500

Epoch 2/40

10/10 [=====] - 3s 299ms/step - loss: 0.6525 - accuracy: 0.5638 - val\_loss: 0.5157 - val\_accuracy: 0.9500

Epoch 3/40

10/10 [=====] - 3s 306ms/step - loss: 0.5928 - accuracy: 0.6383 - val\_loss: 0.3904 - val\_accuracy: 0.9500

Epoch 4/40

10/10 [=====] - 3s 307ms/step - loss: 0.4628 - accuracy: 0.8298 - val\_loss: 0.4797 - val\_accuracy: 0.7500

Epoch 5/40

10/10 [=====] - 3s 288ms/step - loss: 0.3983 - accuracy: 0.8191 - val\_loss: 0.2820 - val\_accuracy: 0.9500

Epoch 6/40

10/10 [=====] - 3s 325ms/step - loss: 0.373  
3 - accuracy: 0.8400 - val\_loss: 0.1934 - val\_accuracy: 1.0000  
Epoch 7/40  
10/10 [=====] - 3s 326ms/step - loss: 0.303  
2 - accuracy: 0.9100 - val\_loss: 0.1718 - val\_accuracy: 1.0000  
Epoch 8/40  
10/10 [=====] - 3s 312ms/step - loss: 0.279  
0 - accuracy: 0.9149 - val\_loss: 0.2136 - val\_accuracy: 0.9500  
Epoch 9/40  
10/10 [=====] - 3s 316ms/step - loss: 0.277  
9 - accuracy: 0.9149 - val\_loss: 0.1722 - val\_accuracy: 0.9500  
Epoch 10/40  
10/10 [=====] - 3s 309ms/step - loss: 0.200  
1 - accuracy: 0.9574 - val\_loss: 0.1962 - val\_accuracy: 0.9500  
Epoch 11/40  
10/10 [=====] - 3s 299ms/step - loss: 0.201  
2 - accuracy: 0.9149 - val\_loss: 0.1182 - val\_accuracy: 1.0000  
Epoch 12/40  
10/10 [=====] - 3s 310ms/step - loss: 0.218  
2 - accuracy: 0.9043 - val\_loss: 0.1165 - val\_accuracy: 1.0000  
Epoch 13/40  
10/10 [=====] - 3s 317ms/step - loss: 0.187  
9 - accuracy: 0.9500 - val\_loss: 0.1677 - val\_accuracy: 1.0000  
Epoch 14/40  
10/10 [=====] - 3s 330ms/step - loss: 0.227  
8 - accuracy: 0.9149 - val\_loss: 0.1260 - val\_accuracy: 0.9500  
Epoch 15/40  
10/10 [=====] - 3s 310ms/step - loss: 0.176  
4 - accuracy: 0.9362 - val\_loss: 0.1776 - val\_accuracy: 0.9000  
Epoch 16/40  
10/10 [=====] - 3s 312ms/step - loss: 0.160  
4 - accuracy: 0.9255 - val\_loss: 0.0842 - val\_accuracy: 1.0000  
Epoch 17/40  
10/10 [=====] - 3s 323ms/step - loss: 0.214  
7 - accuracy: 0.8936 - val\_loss: 0.0938 - val\_accuracy: 0.9500  
Epoch 18/40  
10/10 [=====] - 3s 281ms/step - loss: 0.147  
4 - accuracy: 0.9468 - val\_loss: 0.0638 - val\_accuracy: 1.0000  
Epoch 19/40  
10/10 [=====] - 3s 289ms/step - loss: 0.132  
8 - accuracy: 0.9681 - val\_loss: 0.0636 - val\_accuracy: 1.0000  
Epoch 20/40  
10/10 [=====] - 3s 299ms/step - loss: 0.137  
5 - accuracy: 0.9574 - val\_loss: 0.0799 - val\_accuracy: 1.0000  
Epoch 21/40  
10/10 [=====] - 3s 310ms/step - loss: 0.138  
1 - accuracy: 0.9468 - val\_loss: 0.0828 - val\_accuracy: 1.0000  
Epoch 22/40  
10/10 [=====] - 3s 276ms/step - loss: 0.136  
1 - accuracy: 0.9681 - val\_loss: 0.0671 - val\_accuracy: 1.0000  
Epoch 23/40  
10/10 [=====] - 3s 300ms/step - loss: 0.138  
8 - accuracy: 0.9468 - val\_loss: 0.0762 - val\_accuracy: 0.9500

Epoch 24/40  
10/10 [=====] - 3s 300ms/step - loss: 0.098  
1 - accuracy: 0.9681 - val\_loss: 0.0731 - val\_accuracy: 1.0000  
Epoch 25/40  
10/10 [=====] - 3s 292ms/step - loss: 0.105  
6 - accuracy: 0.9894 - val\_loss: 0.0870 - val\_accuracy: 1.0000  
Epoch 26/40  
10/10 [=====] - 3s 297ms/step - loss: 0.155  
2 - accuracy: 0.9362 - val\_loss: 0.0756 - val\_accuracy: 1.0000  
Epoch 27/40  
10/10 [=====] - 3s 297ms/step - loss: 0.150  
6 - accuracy: 0.9362 - val\_loss: 0.0549 - val\_accuracy: 1.0000  
Epoch 28/40  
10/10 [=====] - 3s 292ms/step - loss: 0.124  
1 - accuracy: 0.9468 - val\_loss: 0.0355 - val\_accuracy: 1.0000  
Epoch 29/40  
10/10 [=====] - 3s 308ms/step - loss: 0.197  
7 - accuracy: 0.9574 - val\_loss: 0.0881 - val\_accuracy: 0.9500  
Epoch 30/40  
10/10 [=====] - 3s 311ms/step - loss: 0.164  
0 - accuracy: 0.9362 - val\_loss: 0.1227 - val\_accuracy: 0.9500  
Epoch 31/40  
10/10 [=====] - 3s 321ms/step - loss: 0.144  
7 - accuracy: 0.9400 - val\_loss: 0.1893 - val\_accuracy: 0.9000  
Epoch 32/40  
10/10 [=====] - 3s 301ms/step - loss: 0.121  
1 - accuracy: 0.9574 - val\_loss: 0.0527 - val\_accuracy: 1.0000  
Epoch 33/40  
10/10 [=====] - 3s 302ms/step - loss: 0.111  
0 - accuracy: 0.9362 - val\_loss: 0.3244 - val\_accuracy: 0.8500  
Epoch 34/40  
10/10 [=====] - 3s 323ms/step - loss: 0.110  
4 - accuracy: 0.9600 - val\_loss: 0.0782 - val\_accuracy: 0.9500  
Epoch 35/40  
10/10 [=====] - 3s 303ms/step - loss: 0.042  
3 - accuracy: 0.9894 - val\_loss: 0.0759 - val\_accuracy: 0.9500  
Epoch 36/40  
10/10 [=====] - 3s 311ms/step - loss: 0.081  
5 - accuracy: 0.9574 - val\_loss: 0.0766 - val\_accuracy: 0.9500  
Epoch 37/40  
10/10 [=====] - 3s 324ms/step - loss: 0.118  
8 - accuracy: 0.9787 - val\_loss: 0.1443 - val\_accuracy: 0.9500  
Epoch 38/40  
10/10 [=====] - 3s 304ms/step - loss: 0.064  
6 - accuracy: 0.9787 - val\_loss: 0.0936 - val\_accuracy: 0.9500  
Epoch 39/40  
10/10 [=====] - 3s 291ms/step - loss: 0.109  
1 - accuracy: 0.9468 - val\_loss: 0.0480 - val\_accuracy: 1.0000  
Epoch 40/40  
10/10 [=====] - 3s 314ms/step - loss: 0.085  
7 - accuracy: 0.9574 - val\_loss: 0.1562 - val\_accuracy: 0.9000

## [5 points] Plot Accuracy and Loss During Training

In [7]:

```
import matplotlib.pyplot as plt

def plot_acc_loss(history, epochs):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']



    plt.figure(figsize=(15, 5))

    plt.subplot(121)
    plt.plot(range(1, epochs), acc[1:], label='Train_acc')
    plt.plot(range(1, epochs), val_acc[1:], label='Test_acc')
    plt.title('Accuracy over ' + str(epochs) + ' Epochs', size=15)
    plt.grid(True)
    plt.legend()

    plt.subplot(122)
    plt.plot(range(1, epochs), loss[1:], label='Train_loss')
    plt.plot(range(1, epochs), val_loss[1:], label='Test_loss')
    plt.title('Loss over ' + str(epochs) + ' Epochs', size=15)
    plt.grid(True)
    plt.legend()

    plt.show()

plot_acc_loss(history, 40)
```



## Plot Test Results



In [8]:

```
import matplotlib.image as mpimg

test_datagen = ImageDataGenerator(rescale=1. / 255)
eval_generator = test_datagen.flow_from_directory(TEST_DIR,target_size=IMAGE_SIZE,
                                                  batch_size=1,shuffle=True,seed
                                                  =42,class_mode="binary")
eval_generator.reset()
pred = model.predict_generator(eval_generator,18,verbose=1)
for index, probability in enumerate(pred):
    image_path = TEST_DIR + "/" + eval_generator.filesnames[index]
    image = mpimg.imread(image_path)
    if image.ndim < 3:
        image = np.reshape(image,(image.shape[0],image.shape[1],1))
        image = np.concatenate([image, image, image], 2)
    # print(image.shape)

    pixels = np.array(image)
    plt.imshow(pixels)


    print(eval_generator.filesnames[index])
    if probability > 0.5:
        plt.title("%.2f" % (probability[0]*100) + "% Normal")
    else:
        plt.title("%.2f" % ((1-probability[0])*100) + "% COVID19 Pneumonia")
    plt.show()
```

Found 18 images belonging to 2 classes.


C:\Users\Li\anaconda3\lib\site-packages\keras\engine\training.py:203  
5: UserWarning: `Model.predict\_generator` is deprecated and will be  
removed in a future version. Please use `Model.predict`, which supports  
generators.

warnings.warn(`Model.predict\_generator` is deprecated and '


18/18 [=====] - 1s 20ms/step  
covid\nejmoa2001191\_f3-PA.jpeg




covid\nejmoa2001191\_f4.jpeg




covid\nejmoa2001191\_f5-PA.jpeg




covid\radiol.2020200490.fig3.jpeg




covid\ryct.2020200028.fig1a.jpeg




covid\ryct.2020200034.fig2.jpeg




covid\ryct.2020200034.fig5-day0.jpeg




covid\ryct.2020200034.fig5-day4.jpeg




covid\ryct.2020200034.fig5-day7.jpeg




normal\NORMAL2-IM-1385-0001.jpeg




normal\NORMAL2-IM-1396-0001.jpeg




normal\NORMAL2-IM-1400-0001.jpeg




normal\NORMAL2-IM-1401-0001.jpeg




normal\NORMAL2-IM-1406-0001.jpeg




normal\NORMAL2-IM-1412-0001.jpeg




normal\NORMAL2-IM-1419-0001.jpeg



normal\NORMAL2-IM-1422-0001.jpeg



normal\NORMAL2-IM-1423-0001.jpeg



## [10 points] TSNE Plot

t-Distributed Stochastic Neighbor Embedding (t-SNE) is a widely used technique for dimensionality reduction that is particularly well suited for the visualization of high-dimensional datasets. After training is complete, extract features from a specific deep layer of your choice, use t-SNE to reduce the dimensionality of your extracted features to 2 dimensions and plot the resulting 2D features.

In [9]:

```
from sklearn.manifold import TSNE

intermediate_layer_model = tf.keras.models.Model(inputs=model.input,
                                                  outputs=model.get_layer('dense').output)
tsne_data_generator = test_datagen.flow_from_directory(DATASET_PATH, target_size=
IMAGE_SIZE,
                                                    batch_size=1, shuffle=False, see
d=42, class_mode="binary")

output = intermediate_layer_model.predict_generator(tsne_data_generator, 130, verb
ose=1)
feature = TSNE(n_components=2).fit_transform(output)
label = tsne_data_generator.classes

covid_x, covid_y, normal_x, normal_y = [], [], [], []

plt.figure()

for index in range(len(feature)):
    if label[index] == 0:
        covid_x.append(feature[index, 0])
        covid_y.append(feature[index, 1])
    else:
        normal_x.append(feature[index, 0])
        normal_y.append(feature[index, 1])

plt.plot(covid_x, covid_y, 'ro', ms=4, label="COVID-19")
plt.plot(normal_x, normal_y, 'bo', ms=4, label="Normal")
plt.legend(loc='upper left')
```



Found 130 images belonging to 2 classes.

1/130 [.....] - ETA: 20s


C:\Users\Li\anaconda3\lib\site-packages\keras\engine\training.py:203  
5: UserWarning: `Model.predict\_generator` is deprecated and will be removed in a future version. Please use `Model.predict`, which supports generators.

warnings.warn(`Model.predict\_generator` is deprecated and '

130/130 [=====] - 8s 58ms/step

Out[9]:

<matplotlib.legend.Legend at 0x26465826308>



## Bonus

In [10]:

```
print('With CPU: \n')  
print('Devices:', tf.config.list_physical_devices())  
print('Time:', total_time)
```

With CPU:

Devices: [PhysicalDevice(name='/physical\_device:CPU:0', device\_type='CPU')]  
Time: 261.01105093955994

In [10]:

```
print('With CPU and GPU: \n')  
  
print('Devices:', tf.config.list_physical_devices())  
  
print('Time:', total_time)
```

With CPU and GPU:

```
Devices: [PhysicalDevice(name='/physical_device:CPU:0', device_type=  
'CPU'), PhysicalDevice(name='/physical_device:GPU:0', device_type='G  
PU')]  
Time: 128.64543533325195
```